
CMPE 493
INTRODUCTION TO
INFORMATION RETRIEVAL

Phrase Queries and
Positional Indexes

Department of Computer Engineering, Boğaziçi University
October 27, 2020

Phrase queries

- ▶ Want to be able to answer queries such as “**Boğaziçi University**” – as a phrase
 - ▶ Thus the sentence “*Sabancı University is at the other side of Boğaziçi Bridge.*” is not a match.
 - ▶ For this, it no longer suffices to store only
 <term : docs> entries
-

A first attempt: Biword indexes

- ▶ Index every consecutive pair of terms in the text as a phrase
- ▶ For example the text “Friends, Romans, Countrymen” would generate the biwords
 - ▶ *friends romans*
 - ▶ *romans countrymen*
- ▶ Each of these biwords is now a dictionary term
- ▶ Two-word phrase query-processing is now immediate.



Longer phrase queries

How would answer the phrase query

boğaziçi university students

boğaziçi university students can be broken into the Boolean query on biwords:

boğaziçi university AND university students

Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase.

Can have false positives!



Extended biwords

- ▶ Parse the indexed text and perform part-of-speech (POS) tagging.
- ▶ Bucket the terms into Nouns (N) and articles/prepositions (X).
- ▶ Call any string of terms of the form NX*N an extended biword.
 - ▶ Each such extended biword is now made a term in the dictionary.
- ▶ Example: ***catcher in the rye***

N X X N
- ▶ Query processing: parse it into N's and X's
 - ▶ Segment query into enhanced biwords
 - ▶ Look up in index: ***catcher rye***

Issues for biword indexes

- ▶ False positives, as noted before
- ▶ Index blowup due to bigger dictionary
 - ▶ Infeasible for more than biwords, big even for them
- ▶ Biword indexes are not the standard solution (for all biwords) but can be part of a compound strategy

Solution 2: Positional indexes

- ▶ In the postings, for each **term** store the position(s) of its tokens in the documents:

<**term**, number of docs containing **term**;

doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>



Positional index example

<**be**: 993427;

1: 7, 18, 33, 72, 86, 231;

2: 3, 149;

4: 17, 191, 291, 430, 434;

5: 363, 367, ...>



Which of docs **1,2,4,5**
could contain "**to be**
or not to be"?

- ▶ For phrase queries, we use a merge algorithm recursively at the document level
- ▶ But we now need to deal with more than just equality



Processing a phrase query

- ▶ Extract inverted index entries for each distinct term: **to**, **be**, **or**, **not**.
 - ▶ Merge their *doc:position* lists to enumerate all positions with “**to be or not to be**”.
 - ▶ **to**:
 - ▶ 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
 - ▶ **be**:
 - ▶ 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
 - ▶ Same general method for proximity searches
-

Proximity queries

- ▶ employment /3 place
 - ▶ Again, here, /k means “within k words of”.
 - ▶ Clearly, positional indexes can be used for such queries; biword indexes cannot.
-

Positional index size

- ▶ Rules of thumb for “English-like” languages:
 - ▶ A positional index is 2–4 as large as a non-positional index
 - ▶ Positional index size 35–50% of volume of original text



Positional index size

- ▶ You can compress position values/offsets
- ▶ Nevertheless, a positional index expands postings storage *substantially*
- ▶ Nevertheless, a positional index is now standardly used because of the power and usefulness of phrase and proximity queries.



Combination schemes

- ▶ These two approaches can be profitably combined
 - ▶ For particular phrases (“**Michael Jackson**”, “**Britney Spears**”) it is inefficient to keep on merging positional postings lists
 - ▶ Even more so for phrases like “**The Who**”

References

- ▶ *Introduction to Information Retrieval*, chapter 2
 - ▶ <http://nlp.stanford.edu/IR-book/information-retrieval-book.html>