
CMPE 493 INTRODUCTION TO INFORMATION RETRIEVAL

Term Weighting and the Vector Space Model

Department of Computer Engineering, Boğaziçi University
November 16, 2020

Ranked retrieval

- ▶ Thus far, our queries have all been Boolean.
 - ▶ Documents either match or don't.
 - ▶ Good for expert users with precise understanding of their needs and the collection.
 - ▶ Also good for applications: Applications can easily consume 1000s of results.
 - ▶ Not good for the majority of users.
 - ▶ Most users incapable of writing Boolean queries (or they are, but they think it's too much work).
 - ▶ Most users don't want to wade through 1000s of results.
 - ▶ This is particularly true of web search.
-

Problem with Boolean search: feast or famine

- ▶ Boolean queries often result in either too few (=0) or too many (1000s) results.
- ▶ Query 1 [Boolean conjunction]:
 - ▶ “*justin bieber istanbul konseri*” → 283,000 hits - **feast**
- ▶ Query 2 [Boolean conjunction]:
 - ▶ “*justin bieber istanbul konseri yeri*” → 0 hits – **famine**
- ▶ It takes a lot of skill to come up with a query that produces a manageable number of hits.
 - ▶ In general: AND gives too few; OR gives too many



Ranked retrieval models

- ▶ Rather than a set of documents satisfying a query expression, in **ranked retrieval models**, the system returns an ordering over the (top) documents in the collection with respect to a query
- ▶ **Free text queries**: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- ▶ In principle, there are two separate choices here, but in practice, ranked retrieval models have normally been associated with free text queries and vice versa



Feast or famine: not a problem in ranked retrieval

- ▶ When a system produces a ranked result set, large result sets are not an issue
 - ▶ Indeed, the size of the result set is not an issue
 - ▶ We just show the top k (≈ 10) results
 - ▶ We don't overwhelm the user
- ▶ Premise: the ranking algorithm works



Importance of ranking:

- ▶ Viewing abstracts: Users are a lot more likely to read the abstracts of the top-ranked pages than the abstracts of the lower ranked pages.
- ▶ Clicking: Distribution is even more skewed for clicking
- ▶ In 1 out of 2 cases, users click on the top-ranked page.
- ▶ Even if the top-ranked page is not relevant, 30% of users will click on it.
- ▶ Getting the ranking right is very important.
- ▶ Getting the top-ranked page right is most important.



Scoring as the basis of ranked retrieval

- ▶ We wish to return in order the documents most likely to be useful to the searcher
- ▶ How can we rank-order the documents in the collection with respect to a query?
- ▶ Assign a score – say in $[0, 1]$ – to each document
- ▶ This score measures how well document and query “match”.



Query-document matching scores

- ▶ We need a way of assigning a score to a query/document pair
- ▶ Let's start with a one-term query
- ▶ If the query term does not occur in the document: score should be 0
- ▶ The more frequent the query term in the document, the higher the score (should be).



Jaccard coefficient

- ▶ Recall from Lecture 5: A commonly used measure of overlap of two sets A and B
 - ▶ $\text{jaccard}(A,B) = |A \cap B| / |A \cup B|$
 - ▶ $\text{jaccard}(A,A) = 1$
 - ▶ $\text{jaccard}(A,B) = 0$ if $A \cap B = \emptyset$
 - ▶ A and B don't have to be the same size.
 - ▶ Always assigns a number between 0 and 1.
-

Jaccard coefficient: Scoring example

- ▶ What is the query-document match score that the Jaccard coefficient computes for each of the two documents below?
 - ▶ Query: *ides of march*
 - ▶ Document 1: *caesar died in march*
 - ▶ Document 2: *of the long march*
-

Issues with Jaccard for scoring

- ▶ It doesn't consider *term frequency* (how many times a term occurs in a document)
- ▶ Rare terms in a collection are more informative than frequent terms. Jaccard doesn't consider this information.



Recall (Lecture 1): Binary term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$



Term-document count matrices

- ▶ Consider the number of occurrences of a term in a document:
 - ▶ Each document is a **count vector** in \mathbb{N}^v : a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0



Bag of words model

- ▶ Vector representation doesn't consider the ordering of words in a document
- ▶ *John is quicker than Mary and Mary is quicker than John have the same vectors*
- ▶ This is called the bag of words model.



Term frequency tf

- ▶ The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- ▶ We want to use tf when computing query-document match scores. But how?
- ▶ Raw term frequency is not what we want:
 - ▶ A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
 - ▶ But not 10 times more relevant.
- ▶ Relevance does not increase proportionally with term frequency.



Log-frequency weighting

- ▶ The log frequency weight of term t in d is

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d}, & \text{if } tf_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

- ▶ $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$, etc.
- ▶ Score for a document-query pair: sum over terms t in both q and d :

$$\text{score} = \sum (1 + \log tf_{t,d})$$

- ▶ The score is 0 if none of the query terms is present in the document.



Document frequency

- ▶ Rare terms are more informative than frequent terms
 - ▶ Recall stop words
- ▶ Consider a term in the query that is rare in the collection (e.g., *arachnophobic*)
- ▶ A document containing this term is very likely to be relevant to the query *arachnophobic*
- ▶ → We want a high weight for rare terms like *arachnophobic*.

Document frequency, continued

- ▶ Frequent terms are less informative than rare terms
- ▶ Consider a query term that is frequent in the collection (e.g., *high, increase, line*)
- ▶ A document containing such a term is more likely to be relevant than a document that doesn't
- ▶ But it's not a sure indicator of relevance.
- ▶ → For frequent terms, we want high positive weights for words like *high, increase, and line*
- ▶ But lower weights than for rare terms.
- ▶ We will use document frequency (df) to capture this.

idf weight

- ▶ df_t is the document frequency of t : the number of documents that contain t
 - ▶ df_t is an inverse measure of the informativeness of t
 - ▶ $df_t \leq N$ (where N is the total number of documents in the collection)
- ▶ We define the idf (inverse document frequency) of t by

$$idf_t = \log_{10}(N/df_t)$$

- ▶ We use $\log(N/df_t)$ instead of N/df_t to “dampen” the effect of idf.



idf example, suppose $N = 1$ million

term	df_t	idf_t
calpurnia	1	6
animal	100	4
sunday	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

$$idf_t = \log_{10}(N/df_t)$$

There is one idf value for each term t in a collection.



Effect of idf on ranking

- ▶ Does idf have an effect on ranking for one-term queries, like
 - ▶ arachnophobic
- ▶ idf has no effect on ranking one term queries
 - ▶ idf affects the ranking of documents for queries with at least two terms
 - ▶ For the query **arachnophobic person**, idf weighting makes occurrences of **arachnophobic** count for more in the final document ranking than occurrences of **person**.



Collection vs. Document frequency

- ▶ The collection frequency of t is the number of occurrences of t in the collection, counting multiple occurrences.
- ▶ Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- ▶ Which word is a better search term (and should get a higher weight)?



tf-idf weighting

- ▶ The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = (1 + \log_{10} \text{tf}_{t,d}) \times \log_{10} (N / \text{df}_t)$$

- ▶ **Best known weighting scheme in information retrieval**
 - ▶ Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - ▶ **Alternative names: tf.idf, tf x idf**
- ▶ Increases with the number of occurrences within a document
- ▶ **Increases with the rarity of the term in the collection**

▶

Final ranking of documents for a query

$$\text{Score}(q, d) = \sum_{t \in q \cap d} \text{tf.idf}_{t,d}$$

▶

Binary → count → weight matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$



Documents as vectors

- ▶ So we have a $|V|$ -dimensional vector space
- ▶ Terms are axes of the space
- ▶ Documents are points or vectors in this space
- ▶ Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- ▶ These are very sparse vectors - most entries are zero.




Queries as vectors

- ▶ [Key idea 1](#): Do the same for queries: represent them as vectors in the space
- ▶ [Key idea 2](#): Rank documents according to their proximity to the query in this space
- ▶ proximity = similarity of vectors
- ▶ proximity \approx inverse of distance
- ▶ Recall: We do this because we want to get away from the you're-either-in-or-out Boolean model.
- ▶ Instead: rank more relevant documents higher than less relevant documents



The Vector-space model




Formalizing vector space proximity

- ▶ First cut: distance between two points
 - ▶ (= distance between the end points of the two vectors)
- ▶ **Euclidean distance?**
- ▶ Euclidean distance is a bad idea ...
- ▶ ... because Euclidean distance is **large** for vectors of **different lengths**.



Why distance is a bad idea



\vec{q} and \vec{d}_2 is large although the distribution of terms in the query q and the distribution of terms in the document d_2 are very similar.



Use angle instead of distance

- ▶ Thought experiment: take a document d and append it to itself. Call this document d' .
- ▶ “Semantically” d and d' have the same content
- ▶ The Euclidean distance between the two documents can be quite large
- ▶ The angle between the two documents is 0, corresponding to maximal similarity.
- ▶ Key idea: Rank documents according to angle with query.




From angles to cosines

- ▶ The following two notions are equivalent.
 - ▶ Rank documents in increasing order of the angle between query and document
 - ▶ Rank documents in decreasing order of $\cos(\text{query}, \text{document})$
- ▶ Cosine is a monotonically decreasing function for the interval $[0^\circ, 180^\circ]$



From angles to cosines



Length normalization

- ▶ A vector can be (length-) normalized by dividing each of its components by its length – for this we use the L_2 norm:

$$\|\vec{x}\|_2 = \sqrt{\sum x_i^2}$$

- ▶ Dividing a vector by its L_2 norm makes it a unit (length) vector (on surface of unit hypersphere)
- ▶ Effect on the two documents d and d' (d appended to itself) from earlier slide: they have identical vectors after length-normalization.
 - ▶ Long and short documents now have comparable weights

cosine(query,document)

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \cdot \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

q_i is the tf-idf weight of term i in the query

d_i is the tf-idf weight of term i in the document

$\cos(\vec{q}, \vec{d})$ is the cosine similarity of \vec{q} and \vec{d} ... or,
equivalently, the cosine of the angle between \vec{q} and \vec{d} .


Cosine for length-normalized vectors

- For length-normalized vectors, cosine similarity is simply the dot product (or scalar product):

$$\cos(\vec{q}, \vec{d}) = \vec{q} \cdot \vec{d} = \sum q_i d_i$$

for q, d length-normalized.

Cosine similarity illustrated



37

Cosine similarity amongst 3 documents

How similar are
the novels

SaS: *Sense and Sensibility* (Jane Austen)

PaP: *Pride and Prejudice* (Jane Austen)

WH: *Wuthering Heights?* (Emily Bronte)

term	SaS	PaP	WH
affection	115	58	20
jealous	10	7	11
gossip	2	0	6
wuthering	0	0	38

Term frequencies (counts)

Note: To simplify this example, we don't do idf weighting.

►

3 documents example contd.

Log frequency weighting

term	SaS	PaP	WH
affection	3.06	2.76	2.30
jealous	2.00	1.85	2.04
gossip	1.30	0	1.78
wuthering	0	0	2.58

After length normalization

term	SaS	PaP	WH
affection	0.789	0.832	0.524
jealous	0.515	0.555	0.465
gossip	0.335	0	0.405
wuthering	0	0	0.588

$\cos(\text{SaS}, \text{PaP}) \approx$

$0.789 \times 0.832 + 0.515 \times 0.555 + 0.335 \times 0.0 + 0.0 \times 0.0$
 ≈ 0.94

$\cos(\text{SaS}, \text{WH}) \approx 0.79$

$\cos(\text{PaP}, \text{WH}) \approx 0.69$



Computing cosine scores

$\text{COSINESCORE}(q)$

```
1  float Scores[N] = 0
2  float Length[N]
3  for each query term t
4  do calculate  $w_{t,q}$  and fetch postings list for t
5      for each pair( $d, \text{tf}_{t,d}$ ) in postings list
6      do  $\text{Scores}[d] += w_{t,d} \times w_{t,q}$ 
7  Read the array Length
8  for each d
9  do  $\text{Scores}[d] = \text{Scores}[d] / \text{Length}[d]$ 
10 return Top K components of Scores[]
```



Summary – vector space ranking

- ▶ Represent the query as a weighted tf-idf vector
- ▶ Represent each document as a weighted tf-idf vector
- ▶ Compute the cosine similarity score for the query vector and each document vector
- ▶ Rank documents with respect to the query by score
- ▶ Return the top K (e.g., $K = 10$) to the user



Computing Scores in a Complete
Search System

Outline

- ▶ Speeding up vector space ranking
- ▶ Putting together a complete search system
 - ▶ Will require learning about a number of miscellaneous topics and heuristics



Efficient cosine ranking

- ▶ Find the K docs in the collection “nearest” to the query $\Rightarrow K$ largest query-doc cosines.
- ▶ Efficient ranking:
 - ▶ Computing a single cosine efficiently.
 - ▶ Choosing the K largest cosine values efficiently.
 - ▶ Can we do this without computing all N cosines?



Special case – unweighted queries

- ▶ No weighting on query terms
 - ▶ Assume each query term occurs only once
- ▶ Then for ranking, don't need to normalize query vector



Faster cosine: unweighted query

FASTCOSINESCORE(q)

```
1  float Scores[N] = 0
2  for each  $d$ 
3  do Initialize  $Length[d]$  to the length of doc  $d$ 
4  for each query term  $t$ 
5  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
6    for each pair( $d, tf_{t,d}$ ) in postings list
7    do add  $wf_{t,d}$  to  $Scores[d]$ 
8  Read the array  $Length[d]$ 
9  for each  $d$ 
10 do Divide  $Scores[d]$  by  $Length[d]$ 
11 return Top  $K$  components of  $Scores[]$ 
```

$w_{t,q}$ set to 1

Figure 7.1 A faster algorithm for vector space scores.




Computing the K largest cosines: selection vs. sorting

- ▶ Typically we want to retrieve the top K docs (in the cosine ranking for the query)
 - ▶ not to totally order all docs in the collection
- ▶ Can we pick off docs with K highest cosines?
- ▶ Let J = number of docs with nonzero cosines
 - ▶ We seek the K best of these J



Use heap for selecting top K

- ▶ Binary tree in which each node's value $>$ the values of children (max-heap)
- ▶ Takes $O(J)$ time to build the heap. Then, for each of K "winners": $O(1)$ time to read the max element, and $O(\log J)$ time to maintain heap property.
- ▶ $O(J)$ time to select top K using heap vs. $O(J \log J)$ time when sorting used.



Bottlenecks

- ▶ Primary computational bottleneck in scoring: cosine computation
- ▶ Can we avoid all this computation?
- ▶ Yes, but may sometimes get it wrong
 - ▶ a doc *not* in the top K may creep into the list of K output docs
 - ▶ Is this such a bad thing?



Cosine similarity is only a proxy

- ▶ User has a task and a query formulation
- ▶ Cosine matches docs to query
- ▶ Thus cosine is anyway a proxy for user happiness
- ▶ If we get a list of K docs “close” to the top K by cosine measure, should be ok



Generic approach

- ▶ Find a set A of *contenders*, with $K < |A| \ll N$
 - ▶ A does not necessarily contain the top K , but has many docs from among the top K
 - ▶ Return the top K docs in A
- ▶ Think of A as pruning non-contenders
- ▶ The same approach is also used for other (non-cosine) scoring functions
- ▶ Will look at several schemes following this approach



Index elimination

- ▶ Basic algorithm FastCosineScore only considers docs containing at least one query term
- ▶ Take this further:
 - ▶ Only consider high-idf query terms
 - ▶ Only consider docs containing many query terms



High-idf query terms only

- ▶ For a query such as *catcher in the rye*
- ▶ **Only accumulate scores from *catcher* and *rye***
- ▶ Intuition: *in* and *the* contribute little to the scores and so don't alter rank-ordering much
- ▶ Benefit:
 - ▶ Postings of low-idf terms have many docs → these (many) docs get eliminated from set A of contenders



Docs containing many query terms

- ▶ Any doc with at least one query term is a candidate for the top K output list
- ▶ **For multi-term queries, only compute scores for docs containing several of the query terms**
 - ▶ Say, at least 3 out of 4
- ▶ **Easy to implement in postings traversal**



3 of 4 query terms

Antony	→	3	4	8	16	32	64	128	
Brutus	→	2	4	8	16	32	64	128	
Caesar	→	1	2	3	5	8	13	21	34
Calpurnia	→	13	16	32					

Scores only computed for docs 8, 16 and 32.

Champion lists

- ▶ Precompute for each dictionary term t , the r docs of highest weight in t 's postings
 - ▶ Call this the champion list for t
 - ▶ (aka fancy list or top docs for t)
- ▶ **Note that r has to be chosen at index build time**
 - ▶ Thus, it's possible that $r < K$
- ▶ At query time, only compute scores for docs in the champion list of some query term
 - ▶ Pick the K top-scoring docs from amongst these

Static quality scores

- ▶ We want top-ranking documents to be both *relevant* and *authoritative*
- ▶ *Relevance is being modeled by cosine scores*
- ▶ *Authority* is typically a query-independent property of a document
- ▶ *Examples of authority signals*
 - ▶ Wikipedia among websites
 - ▶ Articles in certain newspapers
 - ▶ *A paper with many citations*
 - ▶ *(Pagerank)*



Modeling authority

- ▶ Assign a *query-independent* quality score in $[0, 1]$ to each document d
 - ▶ Denote this by $g(d)$
- ▶ *Thus, a quantity like the number of citations is scaled into $[0, 1]$*



Net score

- ▶ Consider a simple total score combining cosine relevance and authority
- ▶ $\text{net-score}(q,d) = g(d) + \text{cosine}(q,d)$
 - ▶ Can use some other linear combination than an equal weighting
- ▶ Now we seek the top K docs by net score



Top K by net score – fast methods

- ▶ First idea: Order all postings by $g(d)$
- ▶ **Key: this is a common ordering for all postings**
- ▶ Thus, can concurrently traverse query terms' postings for
 - ▶ Postings intersection
 - ▶ Cosine score computation



Why order postings by $g(d)$?

- ▶ Under $g(d)$ -ordering, top-scoring docs likely to appear early in postings traversal
- ▶ In time-bound applications (say, we have to return whatever search results we can in 50 ms), this allows us to stop postings traversal early
 - ▶ Short of computing scores for all docs in postings



High and low lists

- ▶ For each term, we maintain two postings lists called *high* and *low*
 - ▶ Think of *high* as the champion list
- ▶ When traversing postings on a query, only traverse *high* lists first
 - ▶ If we get more than K docs, select the top K and stop
 - ▶ Else proceed to get docs from the *low* lists
- ▶ Can be used even for simple cosine scores, without global quality $g(d)$
- ▶ A means for segmenting index into two tiers

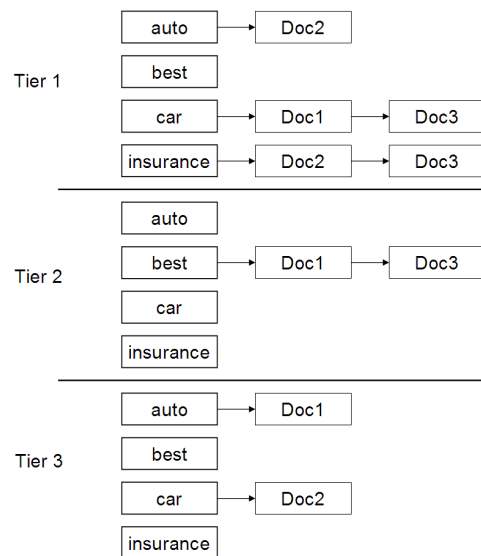


Tiered indexes

- ▶ Break postings up into a hierarchy of lists
 - ▶ Most important
 - ▶ ...
 - ▶ Least important
- ▶ Can be done by $g(d)$ or another measure
- ▶ Inverted index thus broken up into tiers of decreasing importance
- ▶ At query time use top tier unless it fails to yield K docs
 - ▶ If so drop to lower tiers



Example tiered index



Cluster pruning: preprocessing

- ▶ Pick \sqrt{N} docs at random: call these *leaders*
- ▶ For every other doc, pre-compute nearest leader
 - ▶ Docs attached to a leader: its *followers*;
 - ▶ Likely: each leader has $\sim \sqrt{N}$ followers.




Cluster pruning: query processing

- ▶ Process a query as follows:
 - ▶ Given query Q , find its nearest *leader* L .
 - ▶ Seek K nearest docs from among L 's followers.



Visualization



Why use random sampling

- ▶ Fast
- ▶ Leaders reflect data distribution

General variants

- ▶ Have each follower attached to b_1 (e.g., 3) nearest leaders.
- ▶ From query, find b_2 (e.g., 4) nearest leaders and their followers.



Parametric and Zone Indexes



Parametric and zone indexes

- ▶ Thus far, a doc has been a sequence of terms
- ▶ In fact documents have multiple parts, some with special semantics:
 - ▶ Author
 - ▶ Title
 - ▶ Date of publication
 - ▶ Language
 - ▶ Format
 - ▶ etc.
- ▶ These constitute the metadata about a document



Fields

- ▶ We sometimes wish to search by these metadata
 - ▶ E.g., find docs authored by William Shakespeare in the year 1601, containing *alas poor Yorick*
- ▶ Year = 1601 is an example of a field
- ▶ Also, author last name = shakespeare, etc
- ▶ Field or parametric index: postings for each field value
- ▶ Field query typically treated as conjunction
 - ▶ (doc *must* be authored by shakespeare)




Zone

- ▶ A zone is a region of the doc that can contain an arbitrary amount of text e.g.,
 - ▶ Title
 - ▶ Abstract
 - ▶ References ...
- ▶ Build inverted indexes on zones as well to permit querying
- ▶ E.g., “find docs with *merchant* in the title zone and matching the query *gentle rain*”



Example field/zone indexes



Query term proximity

- ▶ Free text queries: just a set of terms typed into the query box – common on the web
 - ▶ Users prefer docs in which query terms occur within close proximity of each other
 - ▶ Let w be the smallest window in a doc containing all query terms, e.g.,
 - ▶ For the query *strained mercy* the smallest window in the doc *The quality of mercy is not strained* is 4 (words)
 - ▶ Would like scoring function to take this into account.
-

Aggregate scores

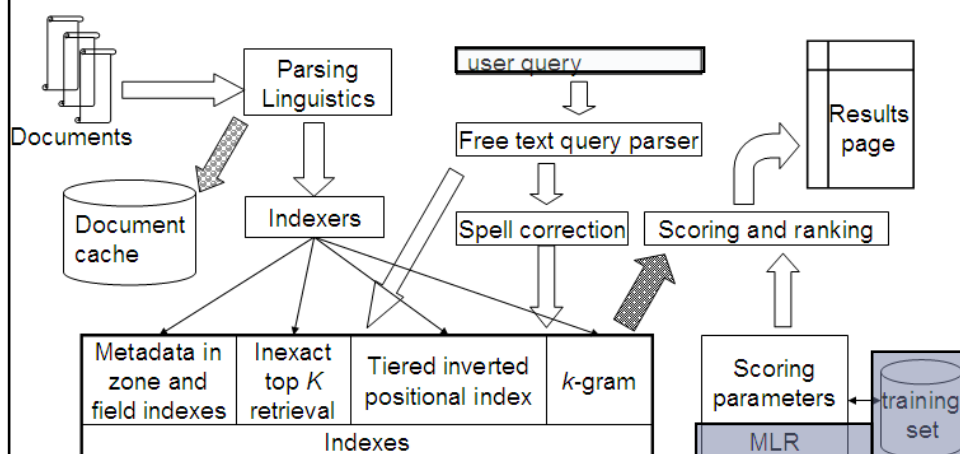
- ▶ We've seen that scoring functions can combine cosine, static quality, proximity, etc.
 - ▶ How do we know the best combination?
 - ▶ Some applications – expert-tuned
 - ▶ Increasingly common: machine-learned
-

Query parsers

- ▶ Free text query from user may in fact spawn one or more queries to the indexes, e.g. query *rising interest rates*
 - ▶ Run the query as a phrase query
 - ▶ If $<K$ docs contain the phrase *rising interest rates*, run the two phrase queries *rising interest* and *interest rates*
 - ▶ If we still have $<K$ docs, run the vector space query *rising interest rates*
 - ▶ Rank matching docs by vector space scoring
- ▶ This sequence is issued by a query parser

▶

Putting it all together



▶

References

- ▶ *Introduction to Information Retrieval*, chapters 6 & 7.
 - ▶ <http://nlp.stanford.edu/IR-book/information-retrieval-book.html>

